

Deep Learning Tutorial

Part II: Advanced concepts

June 16, 2021

Auliya Fitri, Jakob Gawlikowski

Machine Learning Group
Institute of Data Science (DW)
Jena



Curriculum

A: Theoretical introduction – Morning

- I. Introduction and basics
- II. Advanced concepts
- III. Practical application

B. Hands-on seminar – Afternoon

Run prepared Jupyter Notebooks online on Binder, or locally on your own laptop.



Curriculum

II. Advanced concepts

- Regularization
- Convolutional Neural Networks (CNNs)

Inspired by lectures from MIT and Hacettepe Üniversitesi; images taken from these, if not noted otherwise

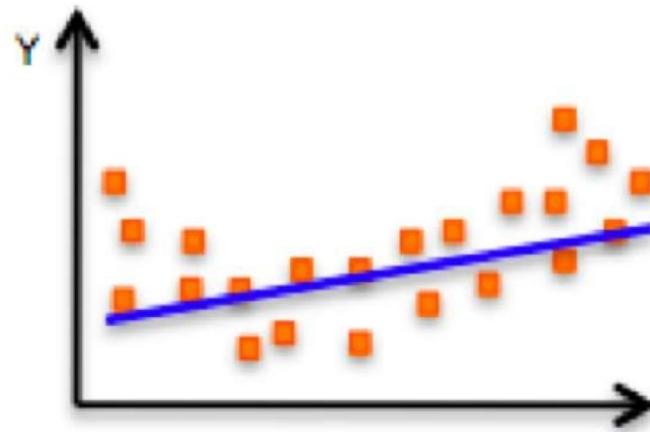


Regularization



Regularization

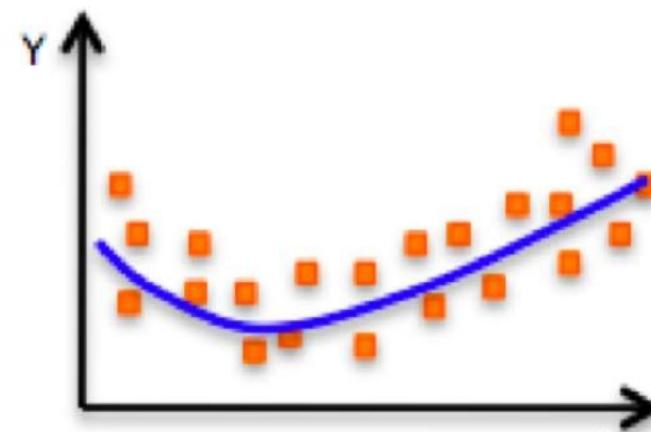
The overfitting problem



Underfitting

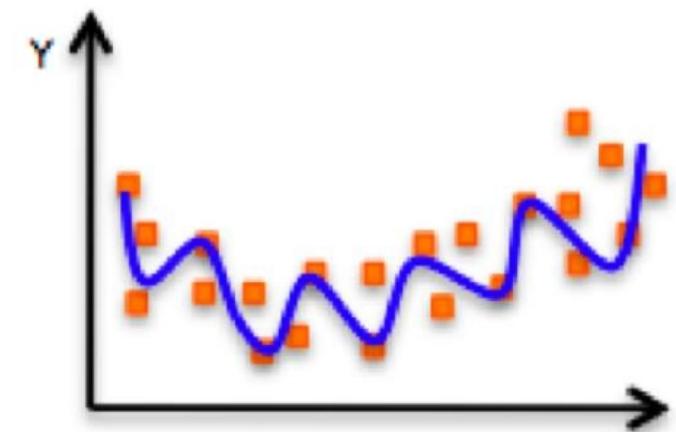
Model does not have capacity
to fully learn the data

= High bias



Ideal fit

Tradeoff



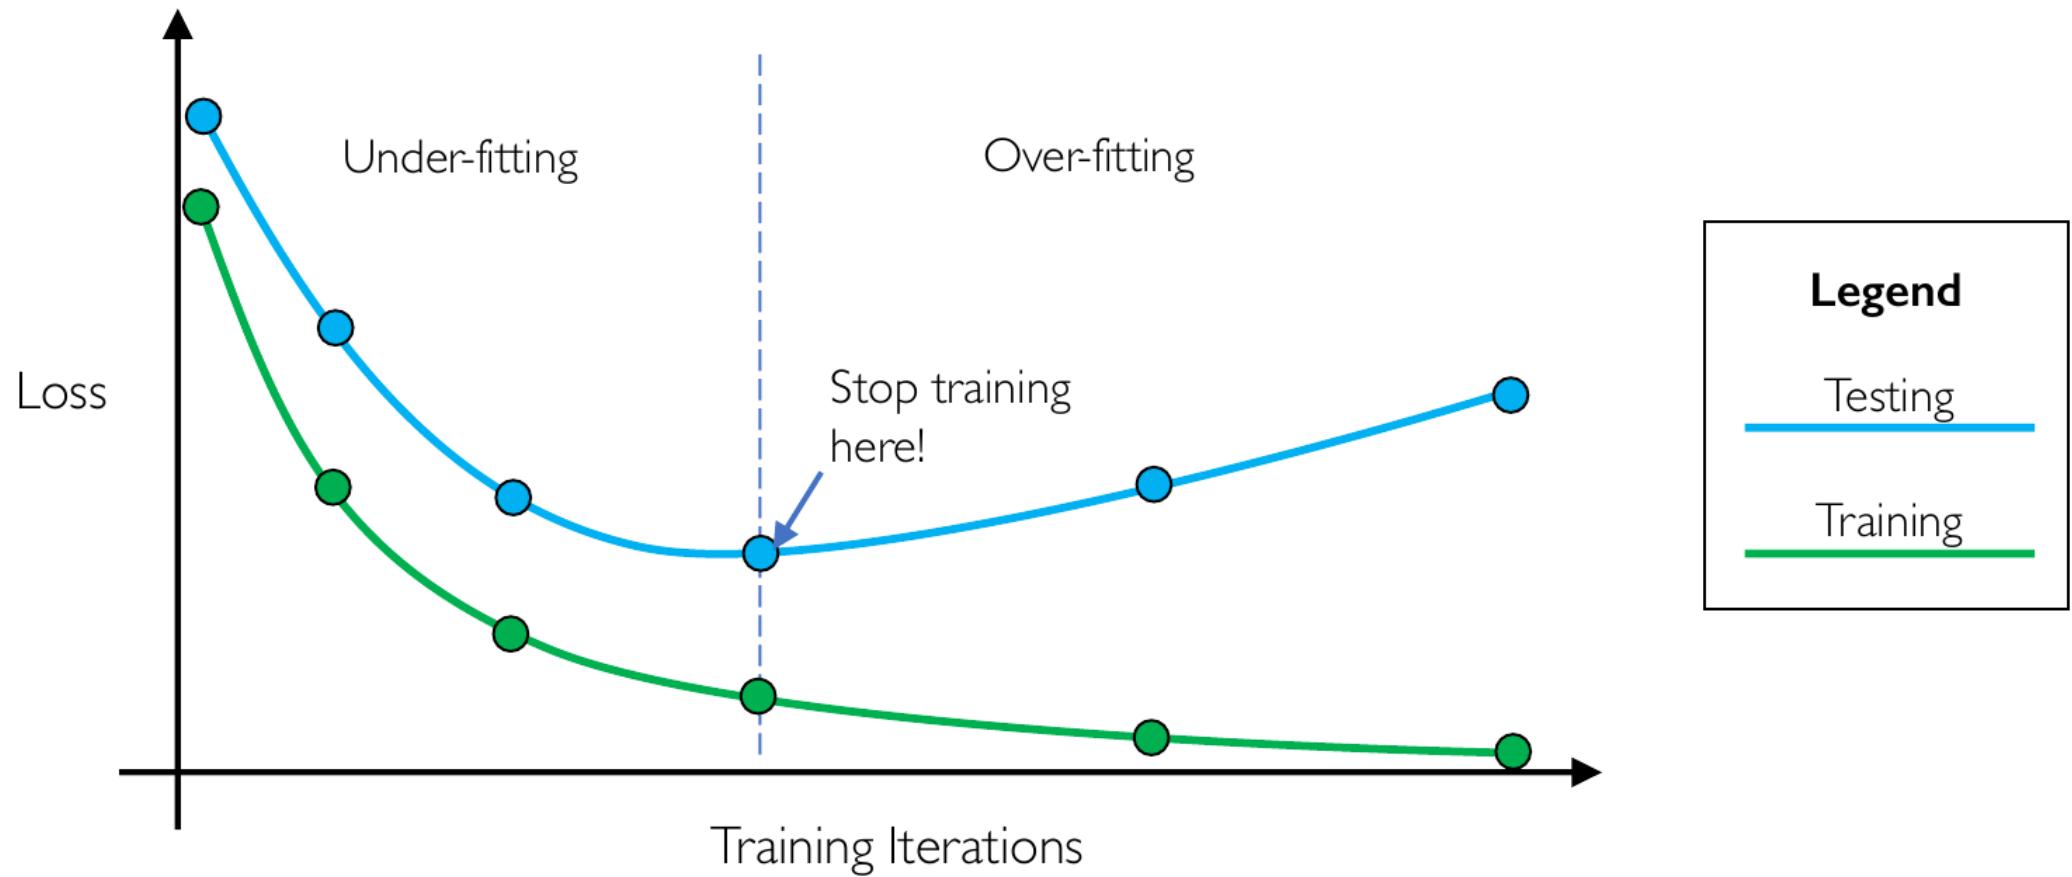
Overfitting

Too complex, extra parameters,
does not generalize well

= High
variance

Regularization

Preventing overfitting: Early stopping



Regularization

Preventing overfitting: L_1 regularization

We want to find the network weights that **achieve the lowest loss**

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)}) + \frac{\lambda}{2} \sum_l \|\mathbf{W}_l\|$$

→ Leads to sparser weights (more zeroes in weights) that are not too adapted to the data at hand



Regularization

Preventing overfitting: Data augmentation

Original



Flip



Random crop



Contrast



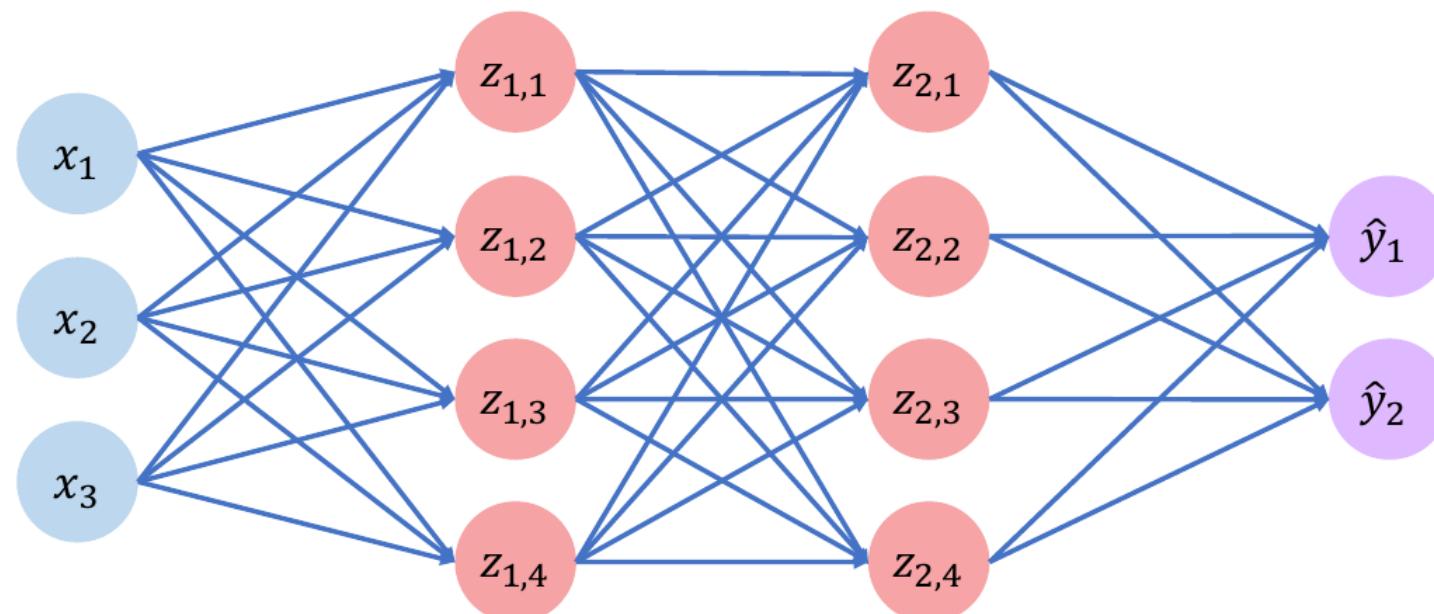
Tint



Regularization

Preventing overfitting: Dropout

- During training, randomly set some activations to 0



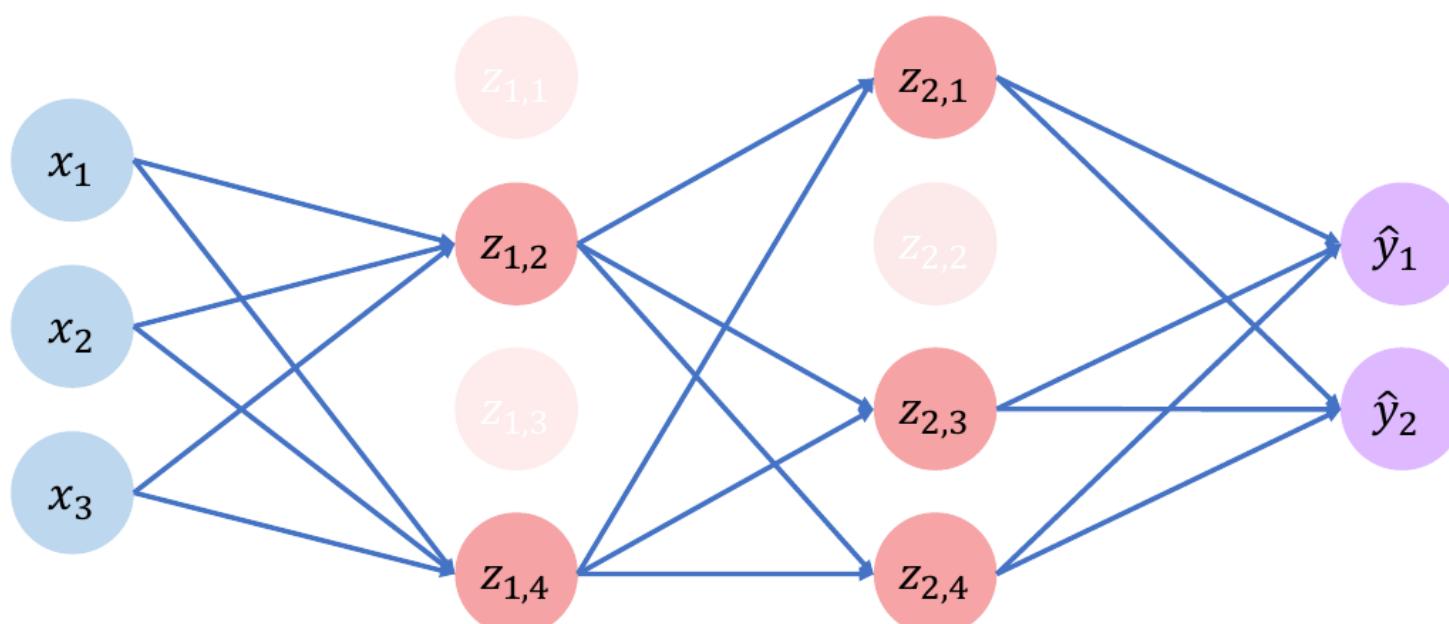
Regularization

Preventing overfitting: Dropout

- During training, randomly set some activations to 0
 - Typically 'drop' 50% of activations in layer
 - Forces network to not rely on any 1 node



tf.keras.layers.Dropout (p=0.5)



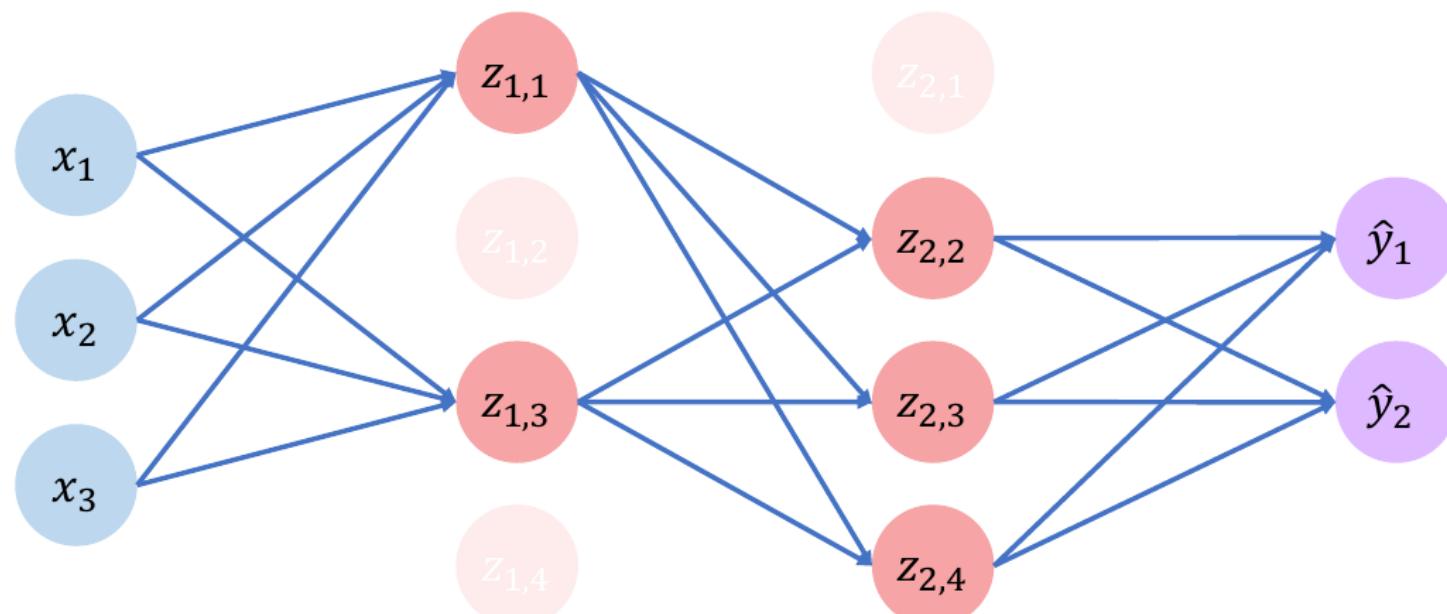
Regularization

Preventing overfitting: Dropout

- During training, randomly set some activations to 0
 - Typically 'drop' 50% of activations in layer
 - Forces network to not rely on any 1 node



tf.keras.layers.Dropout (p=0.5)



Other training considerations

There are a few practical issues (translating to hyper-parameters) for which we do not have enough time, such as:

- Data preprocessing (*normalize your inputs*)
- Network initialization (*use random initializations, e.g. Glorot*)
- Batch normalization (*use*)
- Optimizers, training schedules, learning rates (*when in doubt, use Adam; try different learning rates*)
- Minibatch composition (*shuffle your data sensibly, use large minibatches*)

Fortunately, frameworks and pre-implemented code often already take care of this; nevertheless, you should keep these influences in mind.



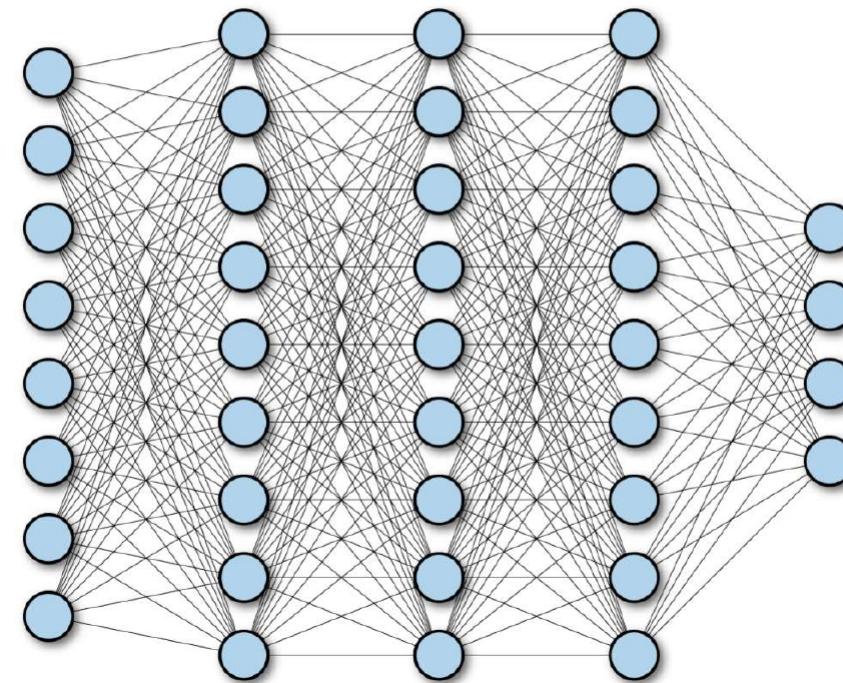
Convolutional Neural Networks (CNNs)



Convolutional Neural Networks

Learning on image data

Fully Connected Neural Network

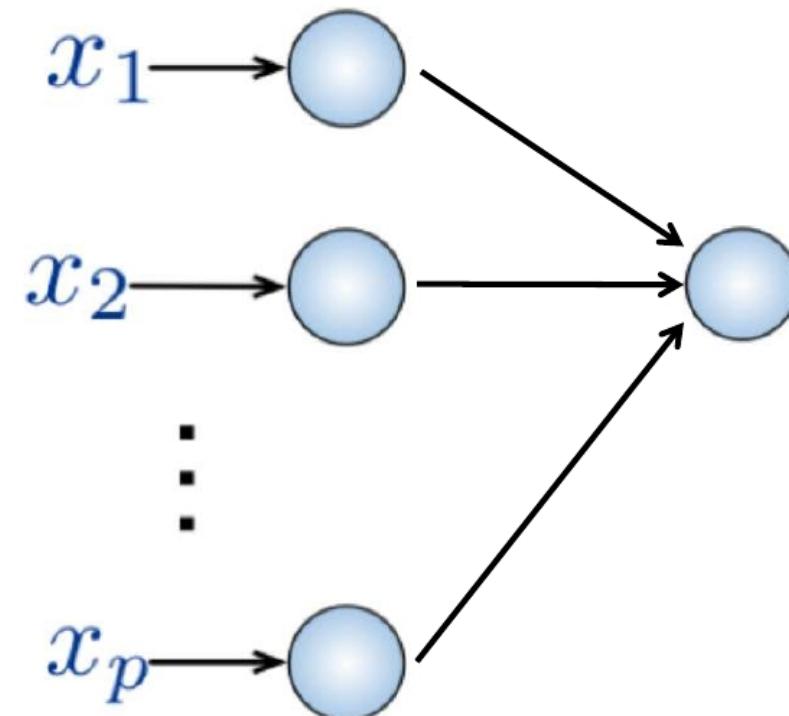


Convolutional Neural Networks

Learning on image data

Input:

- 2D image
- Vector of pixel values



Fully Connected:

- Connect neuron in hidden layer to all neurons in input layer
- No spatial information!
- And many, many parameters!

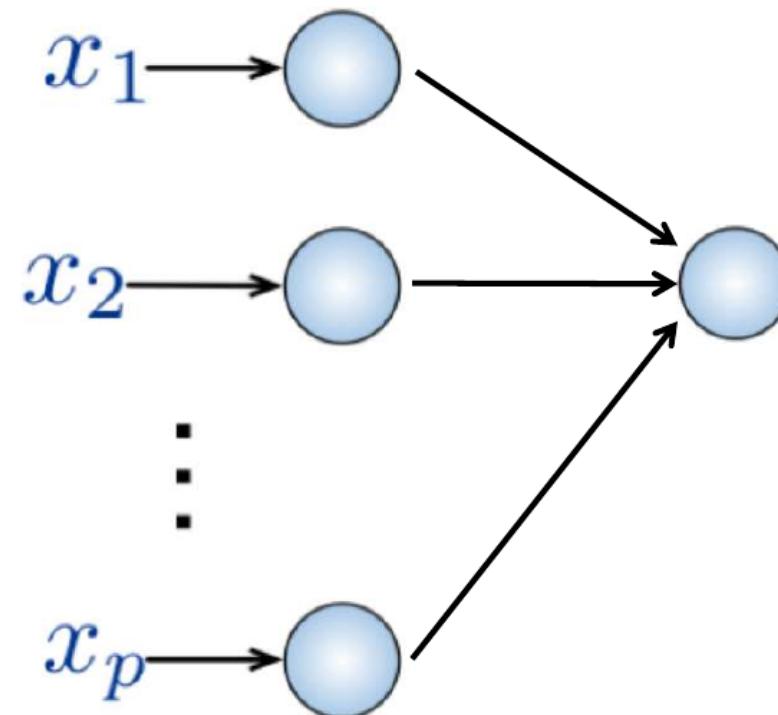


Convolutional Neural Networks

Learning on image data

Input:

- 2D image
- Vector of pixel values



Fully Connected:

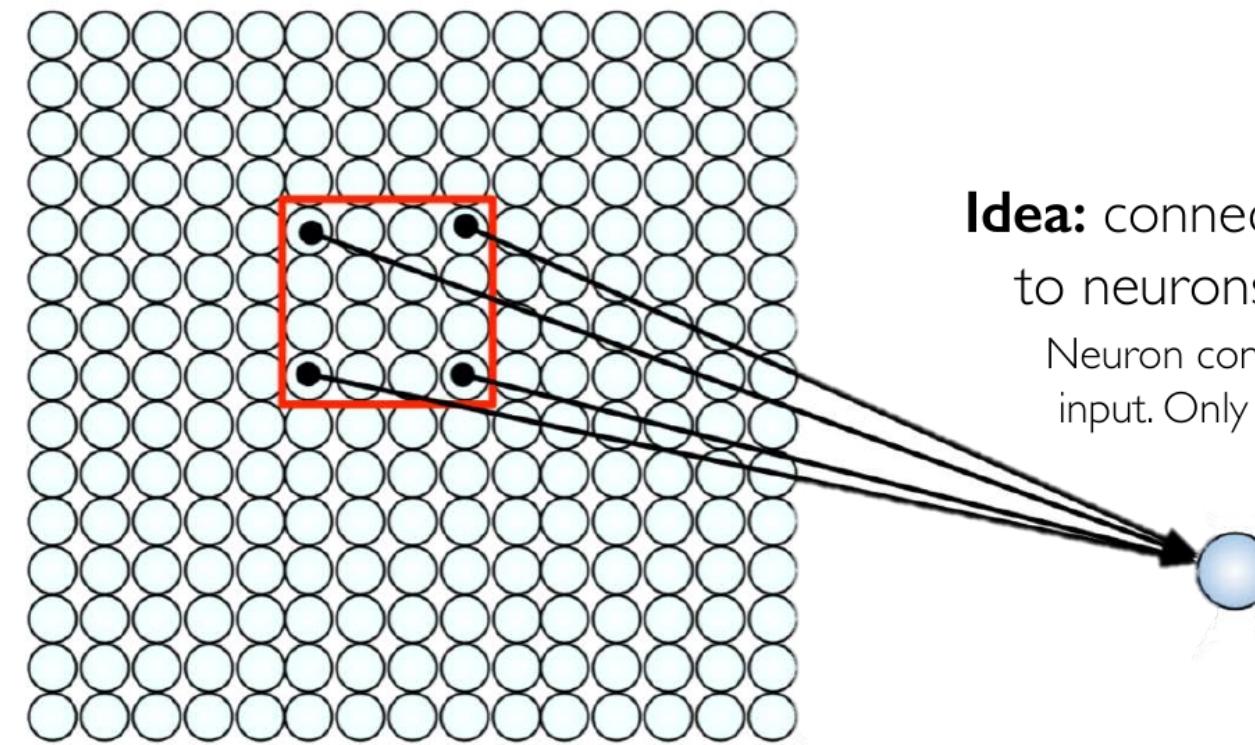
- Connect neuron in hidden layer to all neurons in input layer
- No spatial information!
- And many, many parameters!

How can we use **spatial structure** in the input to inform the architecture of the network?

Convolutional Neural Networks

Learning on image data

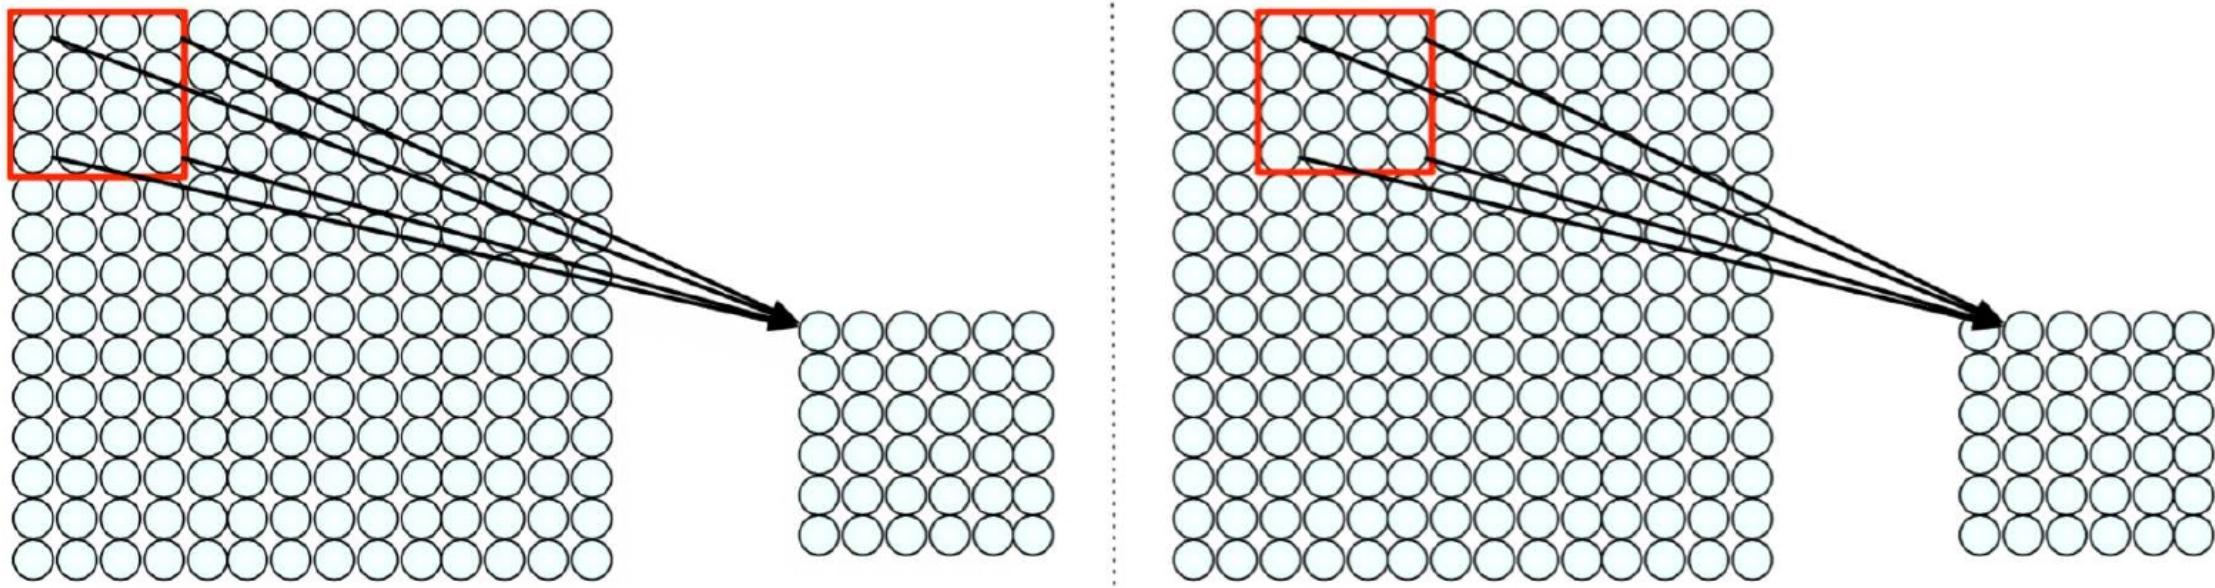
Input: 2D image.
Array of pixel values



Idea: connect patches of input
to neurons in hidden layer.
Neuron connected to region of
input. Only “sees” these values.

Convolutional Neural Networks

Learning on image data



Connect patch in input layer to a single neuron in subsequent layer.

Use a sliding window to define connections.

How can we **weight** the patch to detect particular features?



Convolutional Neural Networks

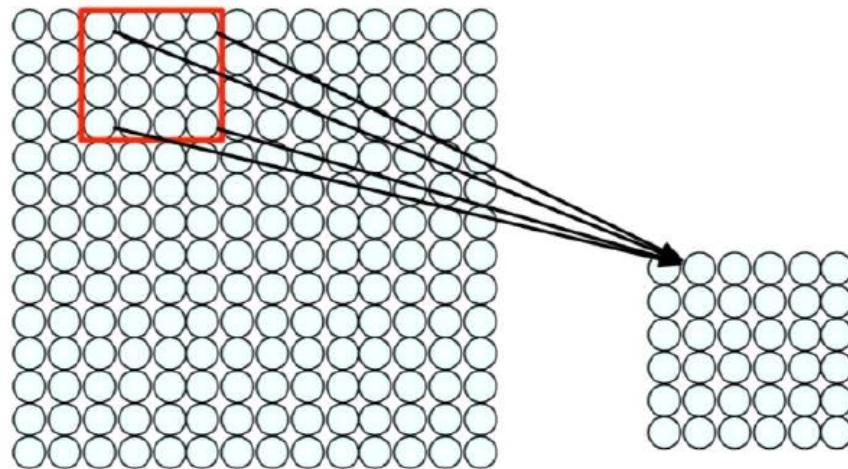
Learning on image data

- 1) Apply a set of weights – a filter – to extract **local features**
- 2) Use **multiple filters** to extract different features
- 3) Spatially **share** parameters of each filter
(features that matter in one part of the input should matter elsewhere)



Convolutional Neural Networks

Learning on image data



- Filter of size 4×4 : 16 different weights
- Apply this same filter to 4×4 patches in input
- Shift by 2 pixels for next patch

This “patchy” operation is **convolution**

- 1) Apply a set of weights – a filter – to extract **local features**
- 2) Use **multiple filters** to extract different features
- 3) **Spatially share** parameters of each filter



Convolutional Neural Networks

Filters to detect features

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	-1	-1	1	-1	-1	-1	-1	
-1	-1	-1	-1	1	-1	-1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	-1	-1	-1	-1	-1	-1	-1	

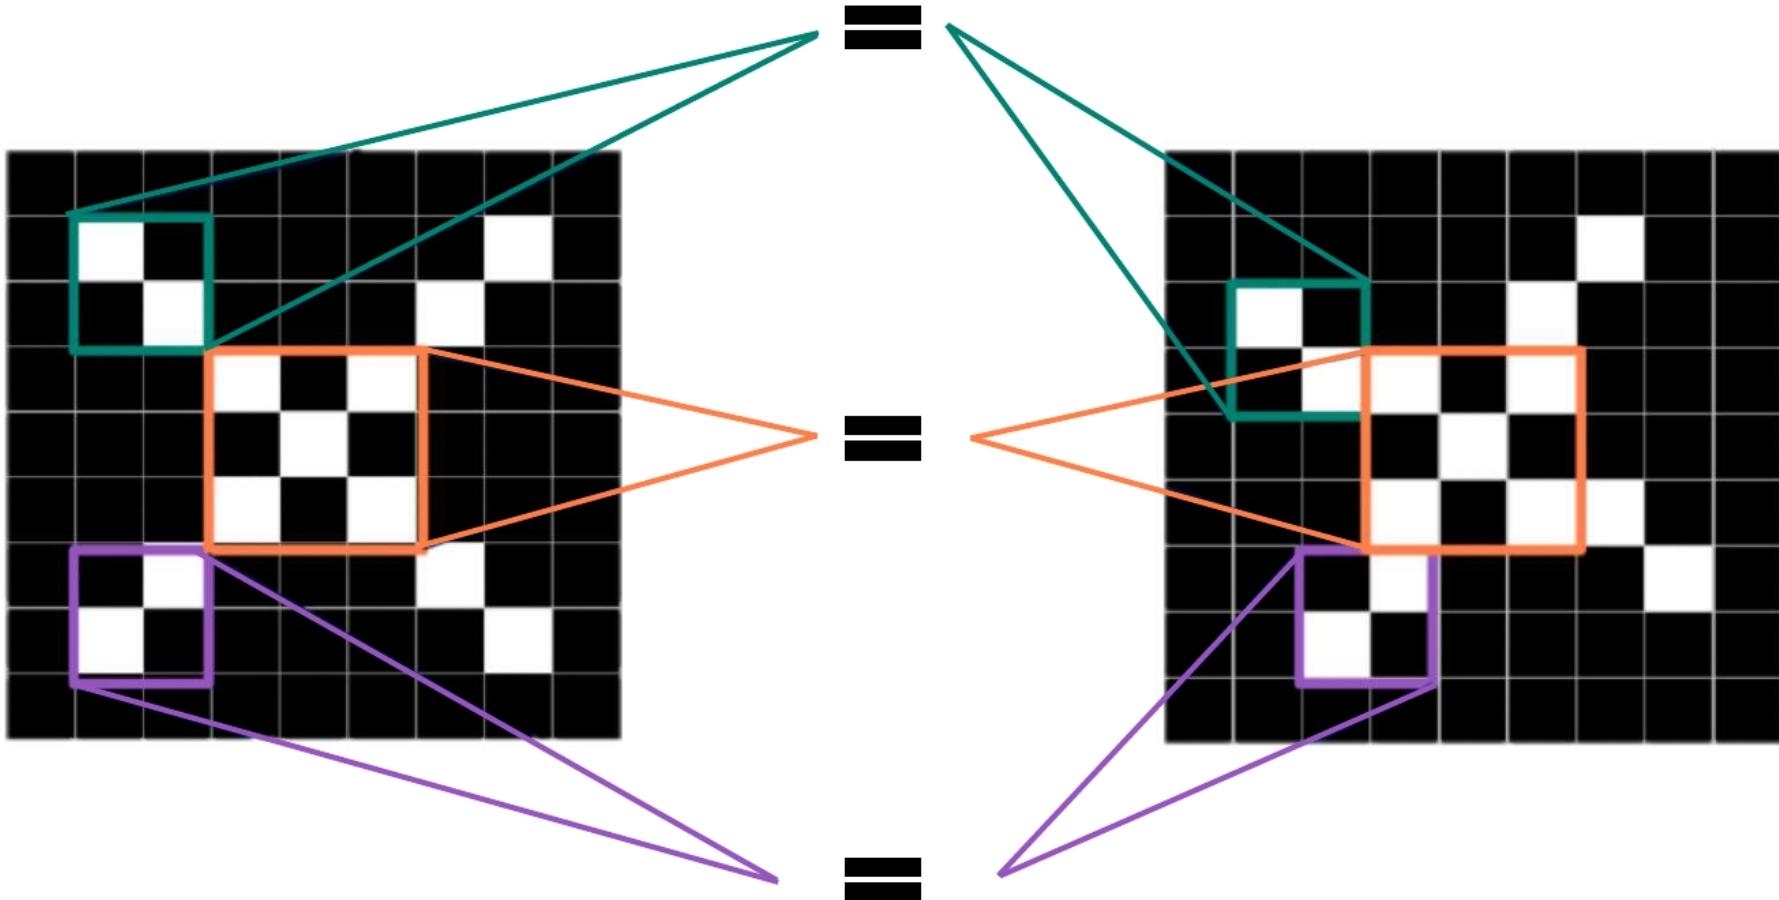


-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

Image is represented as matrix of pixel values... and computers are literal!
We want to be able to classify an X as an X even if it's shifted, shrunk, rotated, deformed.

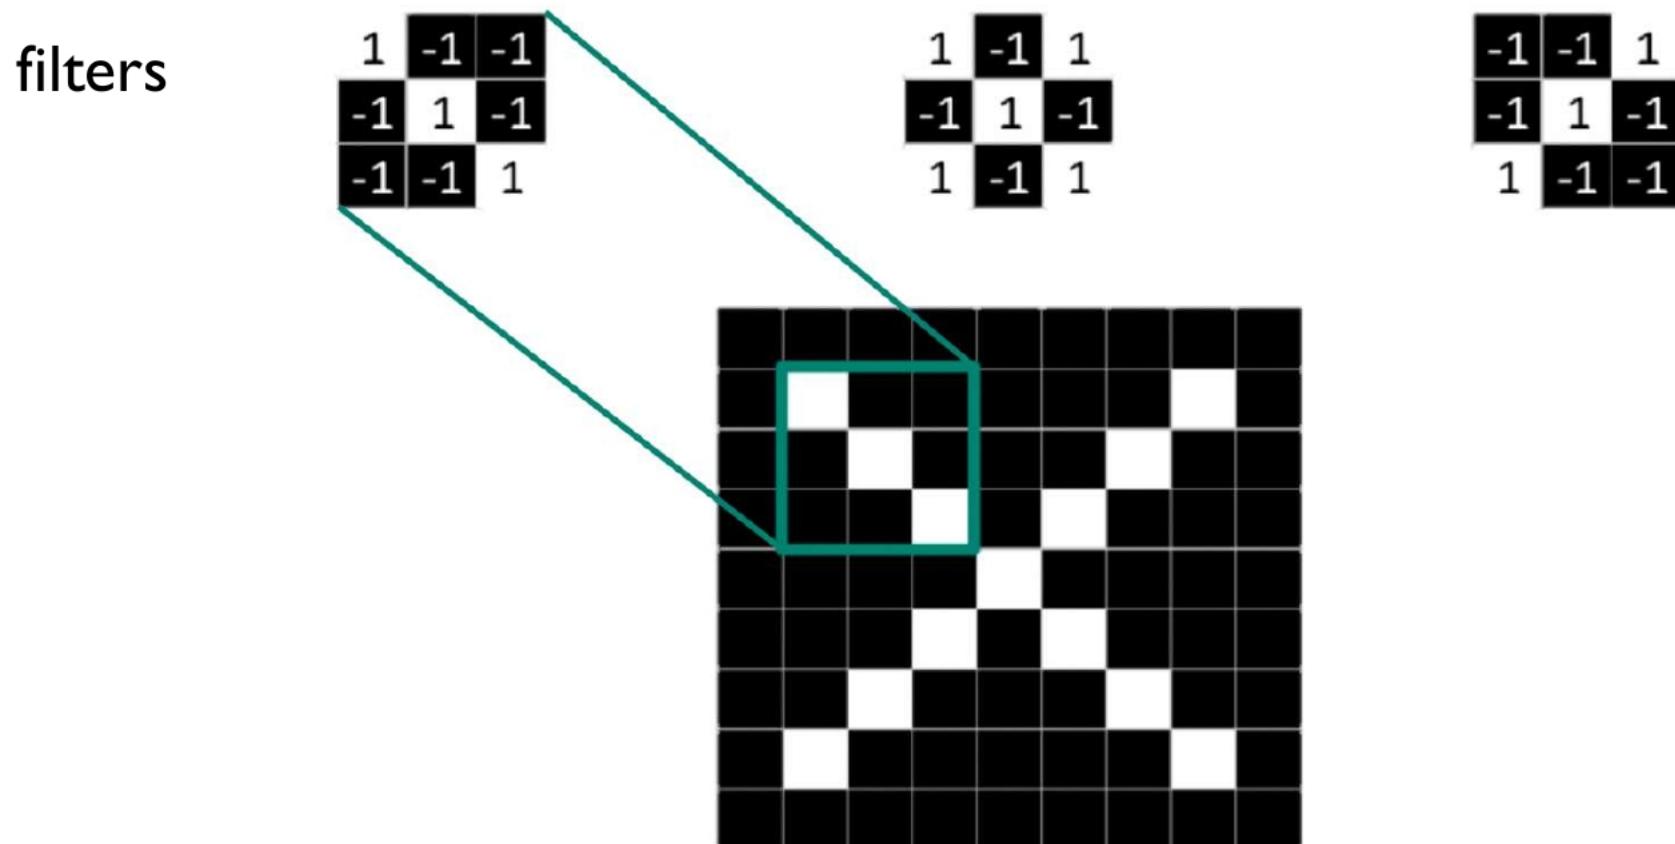
Convolutional Neural Networks

Filters to detect features



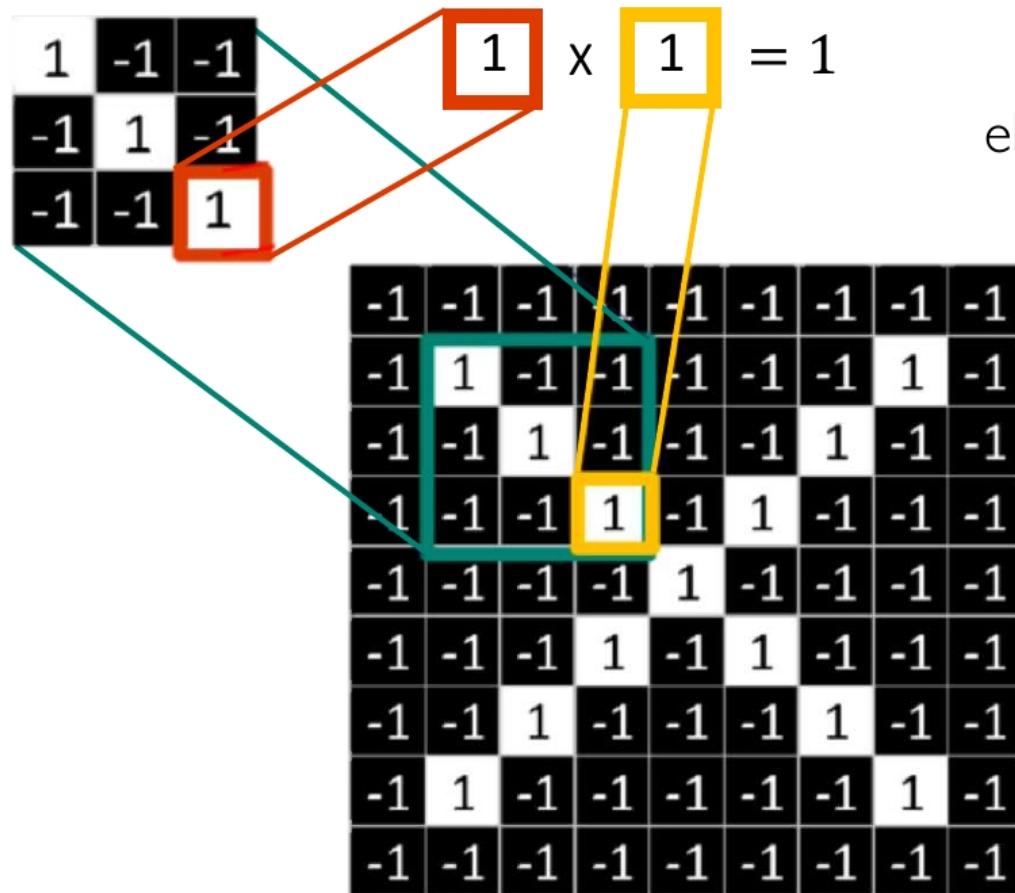
Convolutional Neural Networks

Filters to detect features



Convolutional Neural Networks

The convolution operation



element wise
multiply

add outputs



$$\begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix} = 9$$

Convolutional Neural Networks

The convolution operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0



1	0	1
0	1	0
1	0	1

filter



4		

feature map



Convolutional Neural Networks

The convolution operation

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0



1	0	1
0	1	0
1	0	1

filter

=

4	3	

feature map



Convolutional Neural Networks

The convolution operation

1	1	1	0	0
0	1	1	1	0
0	0	1 _{x1}	1 _{x0}	1 _{x1}
0	0	1 _{x0}	1 _{x1}	0 _{x0}
0	1	1 _{x1}	0 _{x0}	0 _{x1}



1	0	1
0	1	0
1	0	1

filter



4	3	4
2	4	3
2	3	4

feature map



Convolutional Neural Networks

Feature extraction with convolutions



Original



Sharpen



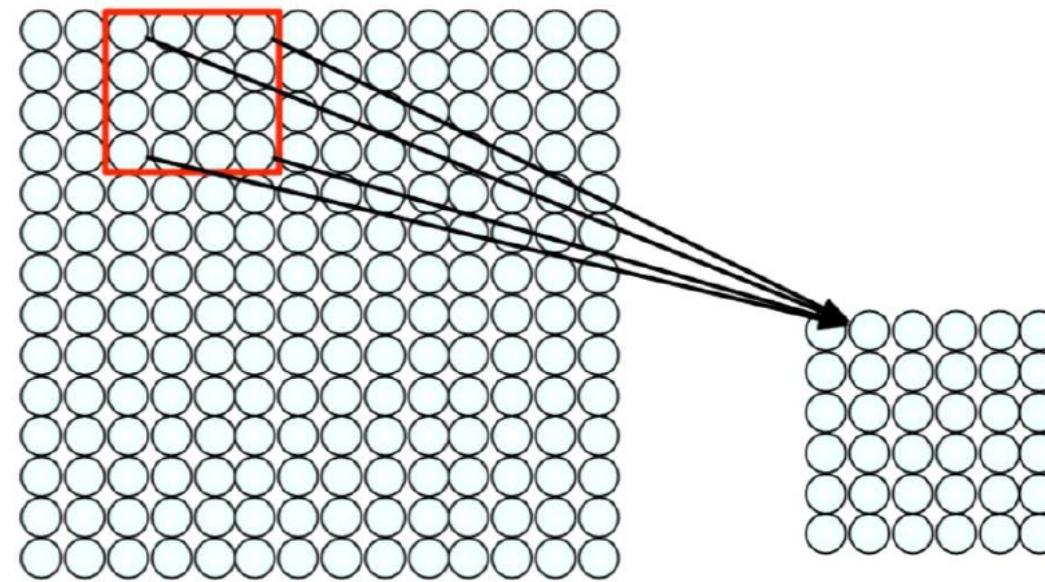
Edge Detect



“Strong” Edge
Detect

Convolutional Neural Networks

Feature extraction with convolutions

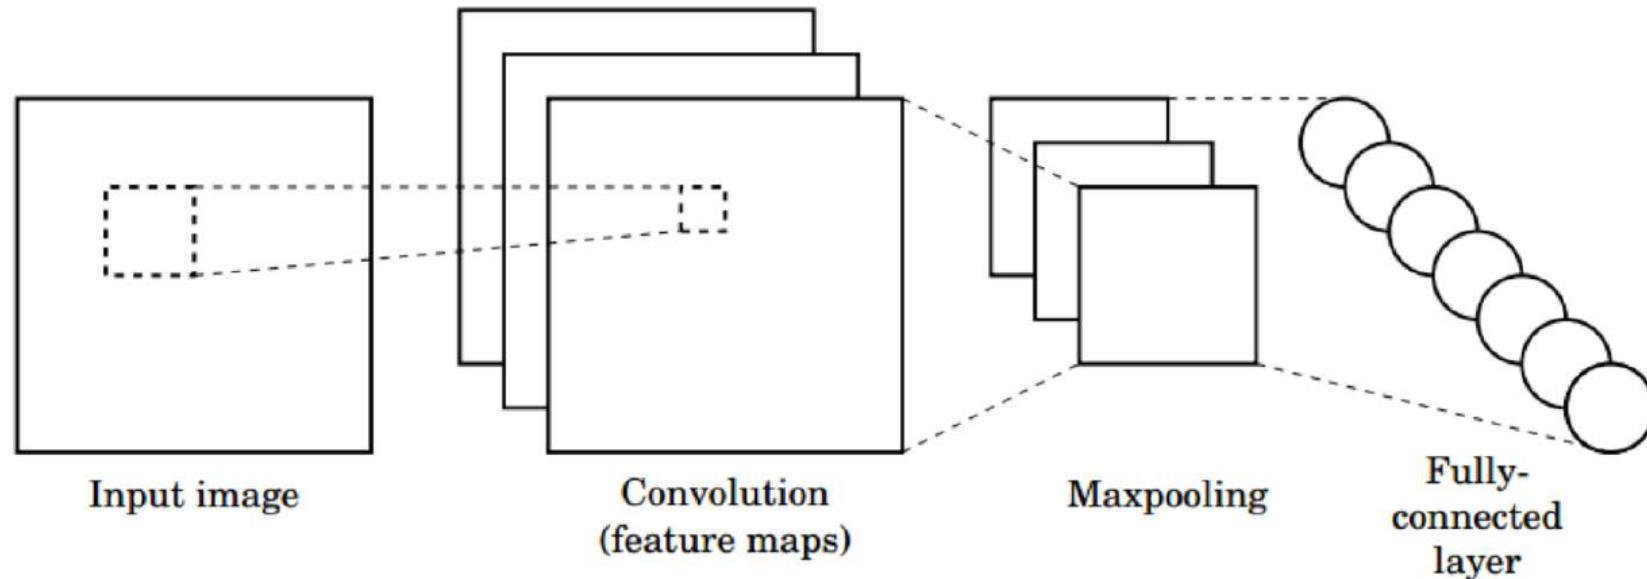


- 1) Apply a set of weights – a filter – to extract **local features**
- 2) Use **multiple filters** to extract different features
- 3) **Spatially share** parameters of each filter



Convolutional Neural Networks

Building a CNN



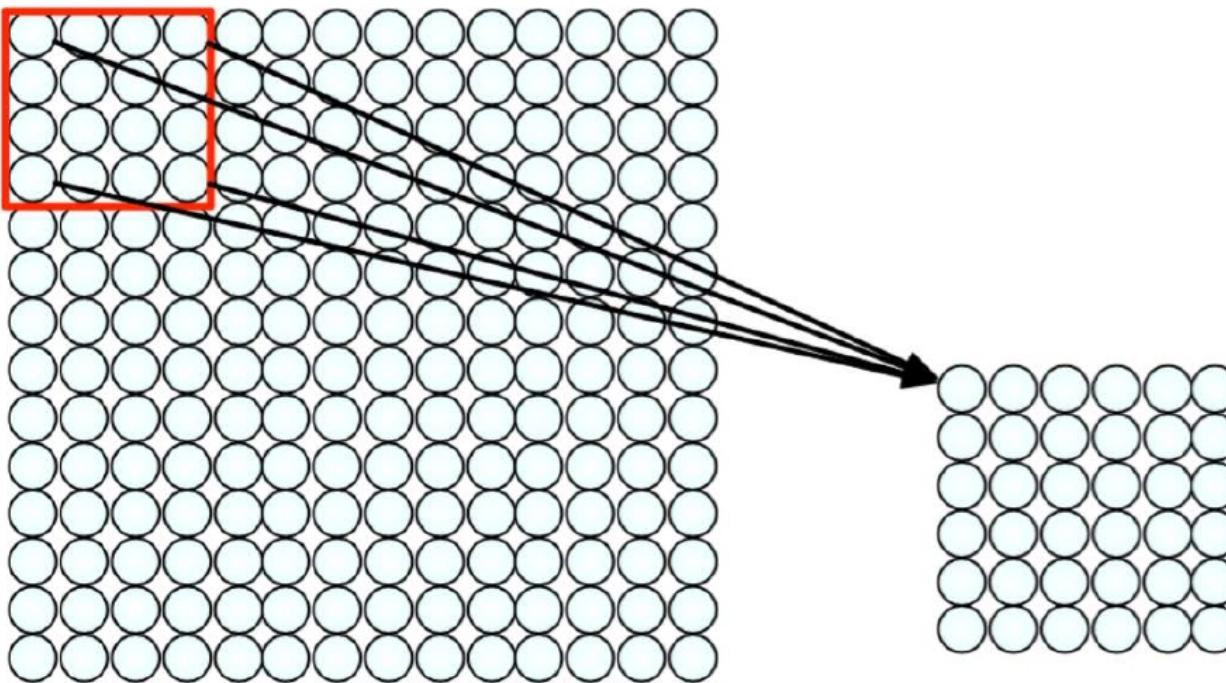
- 1. Convolution:** Apply filters with learned weights to generate feature maps.
- 2. Non-linearity:** Often ReLU.
- 3. Pooling:** Downsampling operation on each feature map.

Train model with image data.

Learn weights of filters in convolutional layers.

Convolutional Neural Networks

Building a CNN



$$\sum_{i=1}^4 \sum_{j=1}^4 w_{ij} x_{i+p,j+q} + b$$

for neuron (p,q) in hidden layer

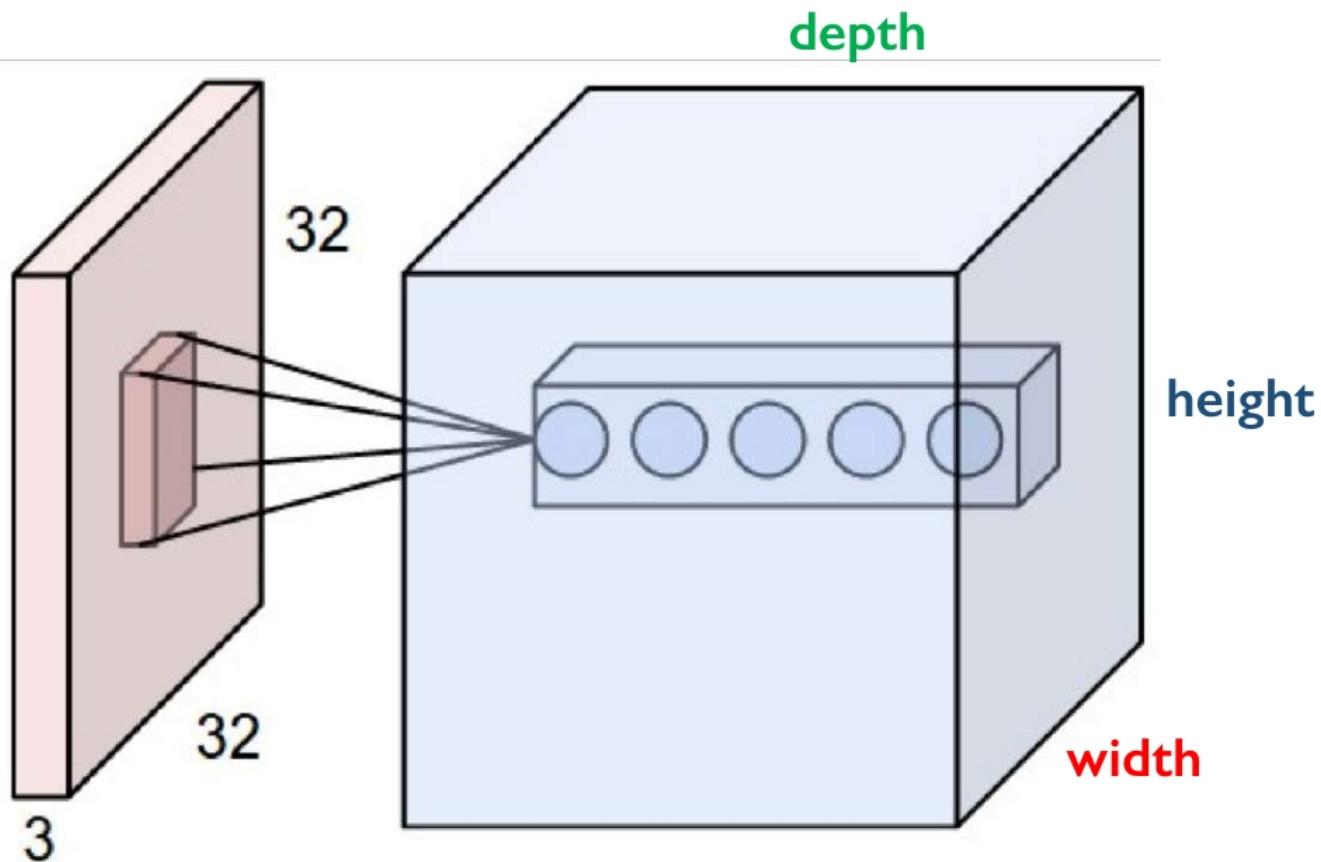
For a neuron in hidden layer:

- Take inputs from patch
- Compute weighted sum
- Apply bias

- 1) applying a window of weights
- 2) computing linear combinations
- 3) activating with non-linear function

Convolutional Neural Networks

Dimensionalities



Layer Dimensions:

$h \times w \times d$

where h and w are spatial dimensions
d (depth) = number of filters

Stride:

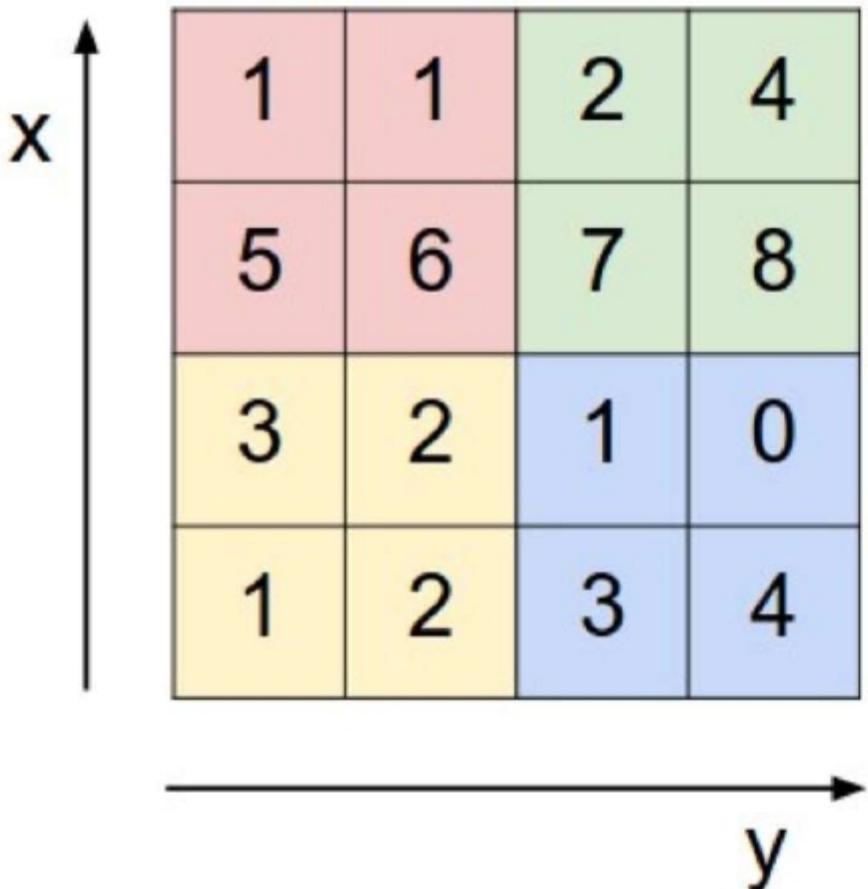
Filter step size

Receptive Field:

Locations in input image that
a node is path connected to

Convolutional Neural Networks

Pooling



max pool with 2x2 filters
and stride 2

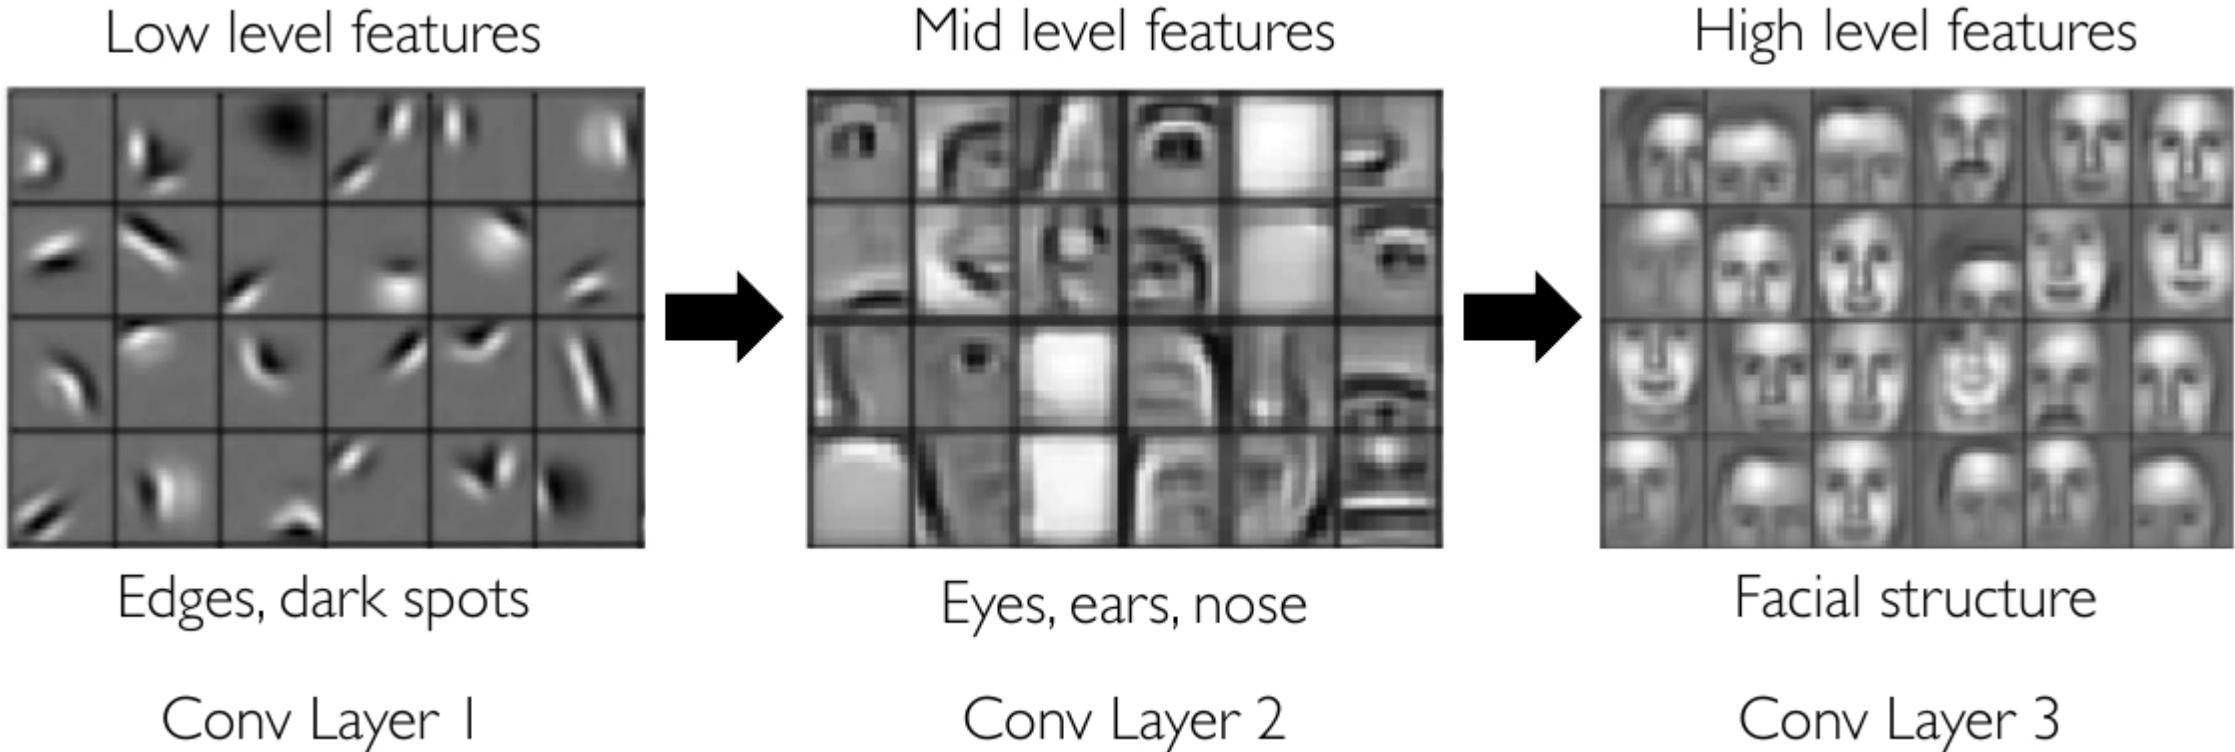


6	8
3	4

- 1) Reduced dimensionality
- 2) Spatial invariance

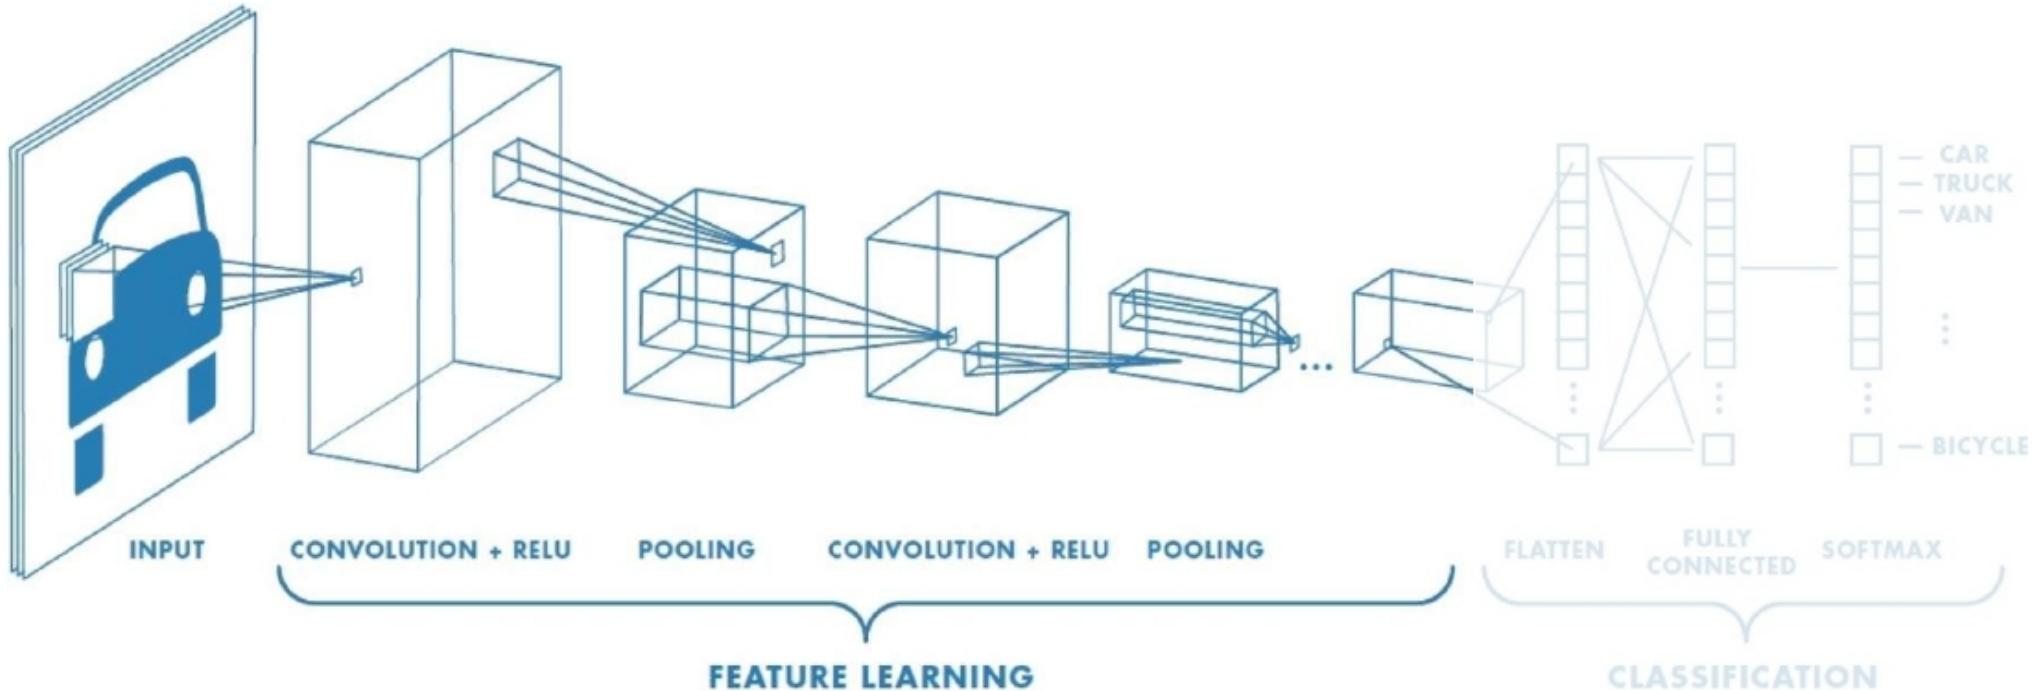
Convolutional Neural Networks

Representation learning



Convolutional Neural Networks

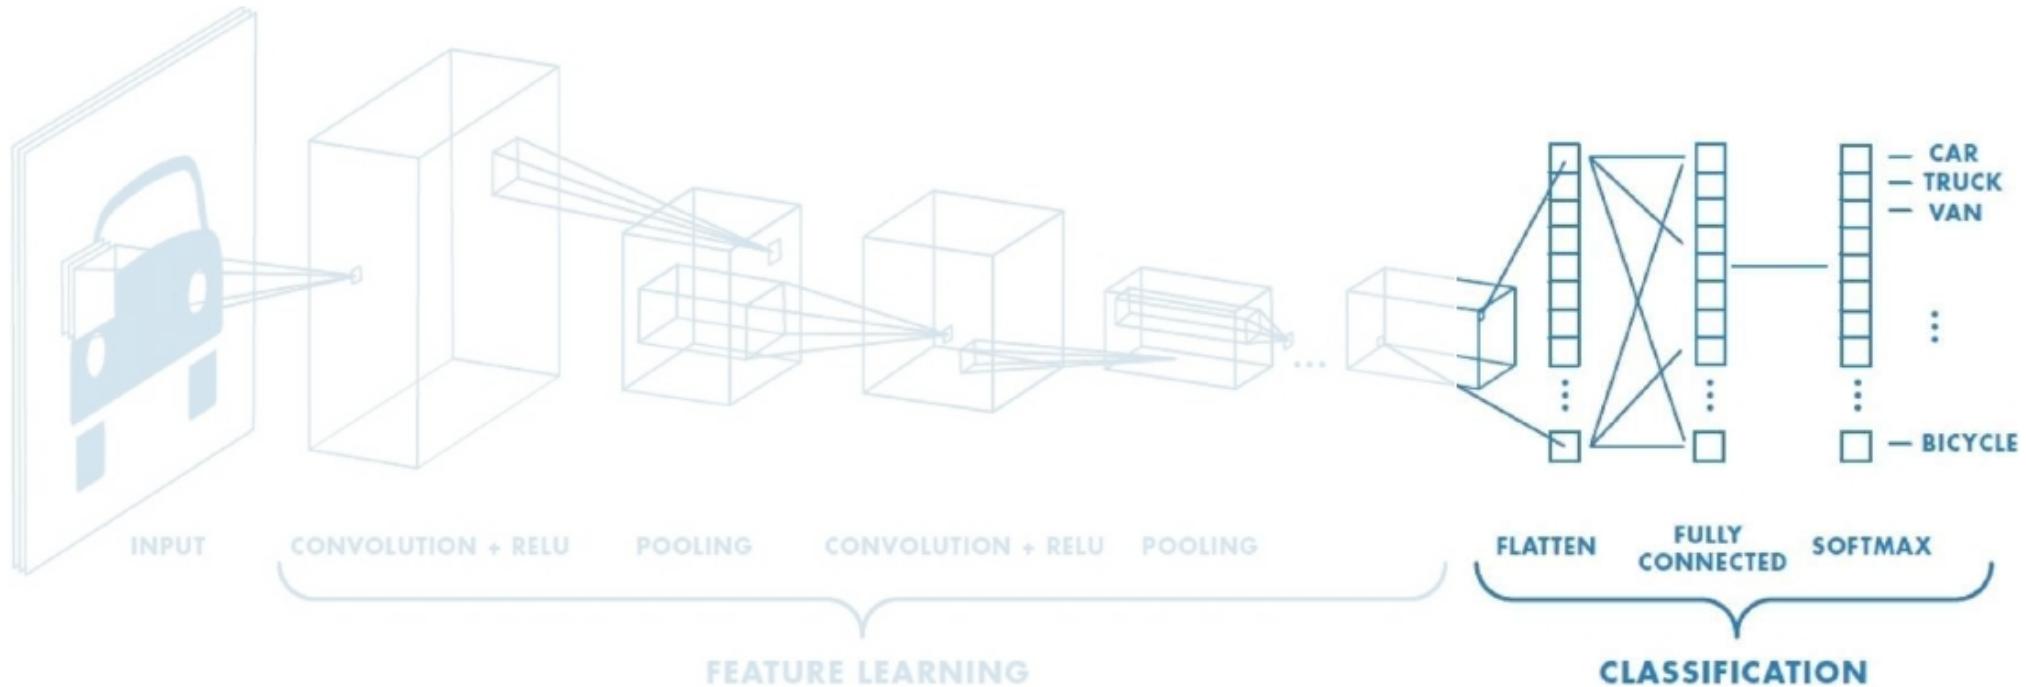
Feature learning



1. Learn features in input image through **convolution**
2. Introduce **non-linearity** through activation function (real-world data is non-linear!)
3. Reduce dimensionality and preserve spatial invariance with **pooling**

Convolutional Neural Networks

Using learned features for classification

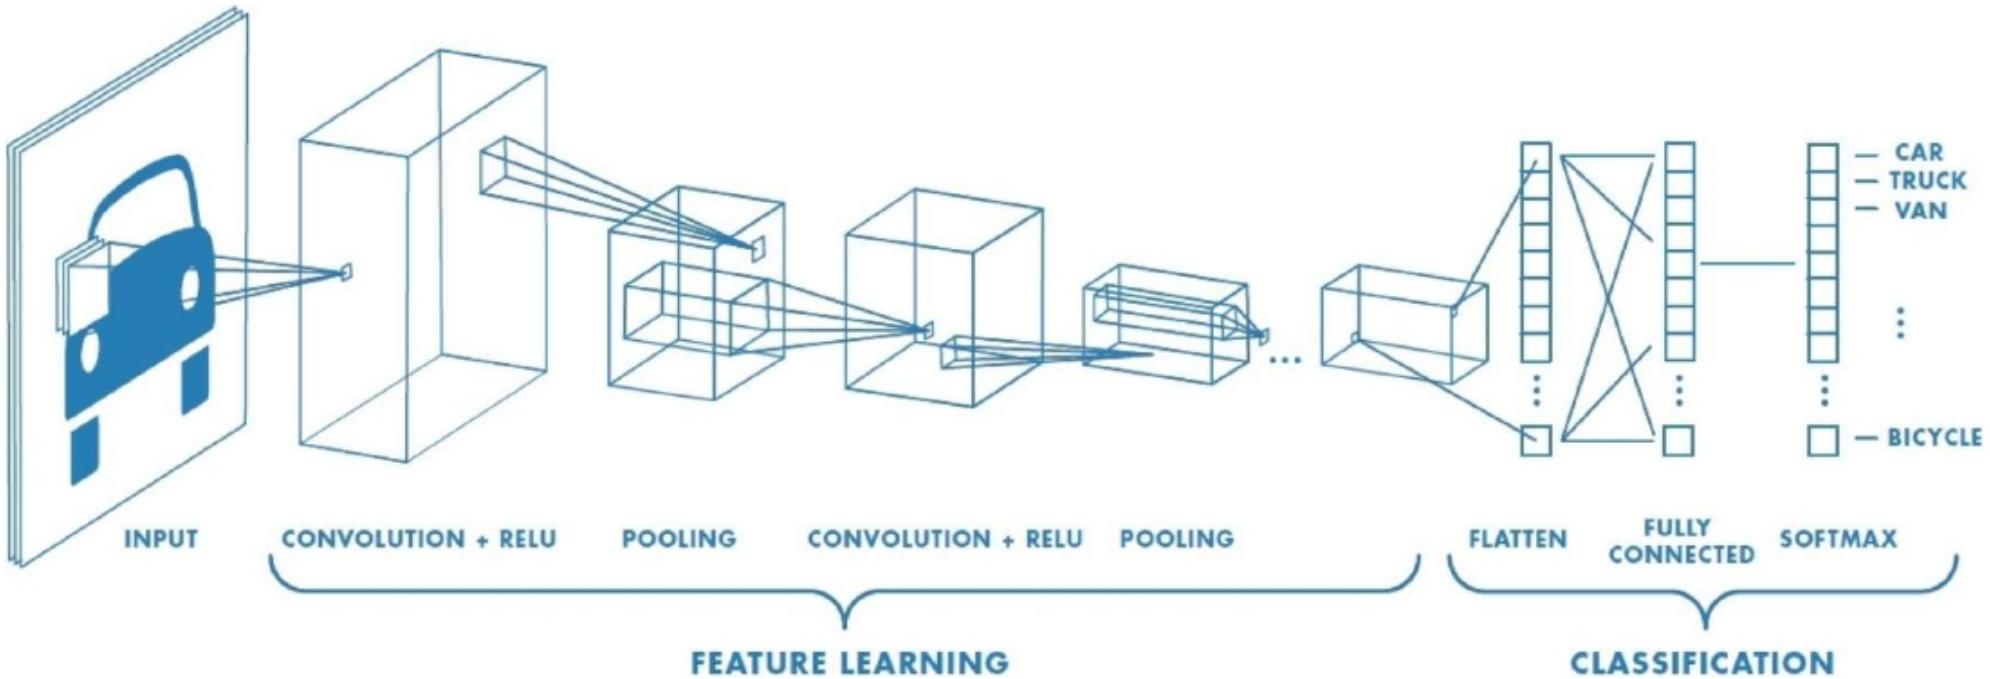


- CONV and POOL layers output high-level features of input
- Fully connected layer uses these features for classifying input image
- Express output as **probability** of image belonging to a particular class

$$\text{softmax}(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

Convolutional Neural Networks

Training with backpropagation



Learn weights for convolutional filters and fully connected layers
Backpropagation: cross-entropy loss

$$J(\theta) = \sum_i y^{(i)} \log(\hat{y}^{(i)})$$

Variants of Neural Networks



Convolution Neural Networks

Image Processing
Computer Vision

Autoencoder

Anomaly Detection
Data Compression
Image Generation

Recurrent Neural Networks

Time Series
Text
Audio



Curriculum

Next: III. Practical application

- Frameworks
- Pre-trained models
- Code & knowledge sources
- (Current research topics)

In B. Hands-on seminar – Try it out:

- Understand PyTorch's Tensor library and neural networks at a high level
- Train a small neural network for regression analysis
- Train a small neural network to classify images