

Introduction to Deep Learning

Part III: Practical application

June 16, 2021

Auliya Fitri, Jakob Gawlikowski

Machine Learning Group
Institute of Data Science (DW)
Jena



Curriculum

A: Theoretical introduction – Morning

- I. Introduction and basics
- II. Advanced concepts
- III. Practical application

B. Hands-on seminar – Afternoon

Run prepared Jupyter Notebooks online on Binder, or locally on your own laptop.



Curriculum

III. Practical application

- Frameworks
- Pre-trained models
- Code & knowledge sources
- (Current research topics)

Inspired by lectures from MIT; images taken from these, if not noted otherwise




Frameworks



Frameworks

What do DL frameworks commonly offer?



Neural network = parametrized, non-linear function

Frameworks

What do DL frameworks commonly offer?

Tensor algebra

X _{a,a}	X _{a,b}	X _{a,c}
X _{b,a}	X _{b,b}	X _{b,c}
X _{c,a}	X _{c,b}	X _{c,c}




Variable(<initial-value>, name=<optional-name>)

```
import tensorflow as tf  
w = tf.Variable(tf.random_normal([3, 3]), name='w')  
y = tf.matmul(x, w)  
relu_out = tf.nn.relu(y)
```

https://www.slideshare.net/tw_dsconf/tensorflow-tutorial

Frameworks


What do DL frameworks commonly offer?



Computation graph: Directed graph of functions, depending on parameters (neuron weights)

Frameworks


What do DL frameworks commonly offer?



Combination of linear (parametrized) and non-linear functions

Frameworks


What do DL frameworks commonly offer?



Not only sequential application of functions

Frameworks

What do DL frameworks commonly offer?




Automatic computation of gradients: all modules are
differentiable!

Frameworks

What do DL frameworks commonly offer?

Automatic
differentiation




<https://towardsdatascience.com/getting-started-with-pytorch-part-1-understanding-how-automatic-differentiation-works-5008282073ec>

Frameworks

What do DL frameworks commonly offer?

Algorithmic optimization,
Parallelization,
Computation on GPU
(CUDA)




https://www.tensorflow.org/guide/data_performance

Frameworks

What do DL frameworks commonly offer?


Pre-defined
components



<https://towardsdatascience.com/how-to-use-batch-normalization-with-tensorflow-and-tf-keras-to-train-deep-neural-networks-faster-60ba4d054b73>

Frameworks

What do DL frameworks commonly offer?



Flexibility for new experiments




Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Non-linear execution
(graphs)



Language usually
Python

Frameworks

Available frameworks



TensorFlow

TensorFlow

- Developed by Google (Brain)
- Fundamentally: Data stream oriented programming
- Used in Gmail, Google Photos, Google Maps, Search, Airbnb, ...
- Easy deployment (including mobile)
- Implemented in Python and C
- Supports Python, C/C++, Java, Go, JavaScript, R, Julia, Swift
- Apache License 2.0
- Version 2.0 is much more user-friendly than 1.0



Frameworks

Available frameworks



PyTorch

- Developed by Facebook
- Based on Torch (Lua)
- “Autograd” module for differentiation
- Dynamic graph – easy debugging
- User-friendly, easy parallelization
- Used by Facebook, Microsoft, Salesforce, ...
- Implemented in Python and C++
- Supports Python, C++
- BSD license



Frameworks

Available frameworks




Keras

- Open source library by François Chollet
- Further abstraction, simplifies many methods and layers
- Can use various backends
- Now integrated in TensorFlow
- Used by Netflix, Uber, Yelp, ...
- Implemented in Python
- Supports Python, R
- MIT license

Frameworks

Available frameworks


Deep Learning Framework Power Scores 2018



<https://keras.io/why-use-keras/>

PyTorch's increasing dominance in research

<https://thegradient.pub/state-of-ml-frameworks-2019-pytorch-dominates-research-tensorflow-dominates-industry/>




Pre-trained models



Pre-trained models

Transfer learning



CNN as universal representations

- First several layers in most CNNs are generic
- They can be reused when training data is comparatively scarce.


Application

- Pre-train on ImageNet classification 1M images
- Cut at some deep conv or FC layer to get features

Pre-trained models

Transfer learning

Deep representations are generic




- A general purpose deep encoder is obtained by chopping off the last layers of a CNN trained on a large dataset.

Pre-trained models

Examples

VGG16





- Proposed in 2014
- Used for image classification
- 16 hidden layers, organized in conv blocks
- Later update: VGG19

<https://neurohive.io/en/popular-networks/vgg16/>

Pre-trained models

Examples



ResNet


- Proposed in 2015
- Used for image classification, object detection, face recognition...
- “Identity shortcut” (or “Highway”) connections
- Stabilizes training of very deep nets
- Many variants

<https://towardsdatascience.com/an-overview-of-resnet-and-its-variants-5281e2f56035>

Pre-trained models


Examples

Inception



(b) Inception module with dimension reductions

- Proposed in 2014
- Used for image classification, object detection, etc.
- Combines filters of various sizes to inception modules – salient features may have different scales
- v1 = GoogLeNet; more versions since then



<https://towardsdatascience.com/a-simple-guide-to-the-versions-of-the-inception-network-7fc52b863202>

Pre-trained models

Examples

More examples from computervision:

- LeNet (1998)
- AlexNet (2012)
- Xception (2016)
- ResNeXt (2017)

Other data types:

- Text embeddings (word or sentence level)
- Audio models pre-trained on Google's Audioset

Availability:

- Many already integrated in Frameworks, e.g. Pytorch
- Or found on project pages, GitHub
- Model architectures can be adapted to other tasks (independent of pre-trained weights)



Code & knowledge sources



Code & knowledge source

Code: Github

- Many free implementations of popular and new methods
- Many conferences now require public code for reproducibility
- Usually a good first step for testing and adapting a new method

keras-yolo3

license MIT

Introduction

A Keras Implementation of YOLOv3 (Tensorflow backend) Inspired by [allanzelener/YAD2K](#).

Quick Start

1. Download YOLOv3 weights from [YOLO website](#).
2. Convert the Darknet YOLO model to a Keras model.
3. Run YOLO detection.

```
wget https://pjreddie.com/media/files/yolov3.weights
python convert.py yolov3.cfg yolov3.weights model_data/yolo.h5
python yolo_video.py [OPTIONS...] --image, for image detection mode, OR
python yolo_video.py [video_path] [output_path (optional)]
```

For Tiny YOLOv3, just do in a similar way, just specify model path and anchor path with `--model model_file` and `--anchors anchor_file`.

Usage

Use `--help` to see usage of `yolo_video.py`:

```
usage: yolo_video.py [-h] [--model MODEL] [--anchors ANCHORS]
                      [--classes CLASSES] [--gpu_num GPU_NUM] [--image]
                      [--input] [--output]
```

Code & knowledge sources

Knowledge: arXiv

- Now often the first point of publication for DL papers (sometimes the only one)
- Not peer-reviewed!
- Helpful tool:

<http://www.arxiv-sanity.com/>

The screenshot shows the arXiv.org interface for the paper "You Only Look Once: Unified, Real-Time Object Detection". The page includes the Cornell University logo, a search bar, and navigation links for "All fields" and "Search". The main content area displays the paper's title, authors, submission date, and a detailed abstract. Below the abstract are sections for "Subjects", "Bibliographic data", and "Submission history". The right sidebar provides download options (PDF, Other formats), browse context (cs.CV), and links to references, citations, and related resources.

Code & knowledge sources

Knowledge: Blogs

The Tensor

The Loss Function

The Autograd

The nn Module

The Autograd

Machine Learning

Teaching the learners.

Following

Follow to get great stories about Machine Learning in your inbox and on your homepage

RELATED TOPICS

- ARTIFICIAL INTELLIGENCE
- DATA SCIENCE
- PROGRAMMING
- TECHNOLOGY
- CYBERSECURITY

POPULAR IN MACHINE LEARNING

- Top 3 Pandas Functions You Don't Know About (Probably)
- A Comprehensive Guide to Convolutional Neural Networks—the ELSI way
- Data Science Books you should read in 2020
- Machine Learning Basics with the K-Nearest Neighbors Algorithm

LATEST

Sound classification using Images

Learn how to classify audio files using spectrograms

Dipam Vasani in Towards Data Science

May 14 · 5 min read ★

<https://medium.com/topic/artificial-intelligence>
<https://medium.com/topic/machine-learning>

Real-time Object Detection with YOLO, YOLOv2 and now YOLOv3



Jonathan Hui

Follow

Mar 18, 2018 · 18 min read




You only look once (YOLO) is an object detection system targeted for real-time processing. We will introduce YOLO, YOLOv2 and YOLO9000 in this article. For those only interested in YOLOv3, please forward to the bottom of the article. Here is the accuracy and speed comparison provided by the YOLO web site.

Model	Train	Test	mAP	FLOPS	FPS
SSD300	COCO trainval	test-dev	41.2	-	46
SSD500	COCO trainval	test-dev	46.5	-	19
YOLOv2 608x608	COCO trainval	test-dev	48.1	62.94 Bn	40
Tiny YOLO	COCO trainval	-	-	7.07 Bn	200
SSD321	COCO trainval	test-dev	45.4	-	16
DSSD321	COCO trainval	test-dev	46.1	-	12
R-FCN	COCO trainval	test-dev	51.9	-	12
SSD513	COCO trainval	test-dev	50.4	-	8
DSSD513	COCO trainval	test-dev	53.3	-	6
FPN FRCN	COCO trainval	test-dev	59.1	-	6
Retinanet 50-500	COCO trainval	test-dev	59.0	-	14

Code & knowledge sources

Knowledge: Conferences (selection)



...and of course many
domain-specific
conferences






Association for
Computational
Linguistics



Code & knowledge sources

Knowledge: Books



[http://
www.deeplearningbook.org/](http://www.deeplearningbook.org/)

Code & knowledge sources

Knowledge: Online courses

Browse > Data Science > Machine Learning

Offered By

Stanford

Machine Learning

★★★★★ 4.9 120,193 ratings • 29,510 reviews

Enroll for Free
Started Nov 25

Financial aid available

2,655,317 already enrolled!

About Syllabus Reviews Instructors Enrollment Options FAQ

About this Course

7,781,080 recent views

Machine learning is the science of getting computers to act without being explicitly programmed. In the past decade, machine learning has given us self-driving cars, practical speech recognition, effective web search, and a vastly improved understanding of the human genome. Machine learning is so pervasive today that you probably use it dozens of times a day without knowing it. Many researchers also think it is the best way to make

[SHOW ALL](#)

SKILLS YOU WILL GAIN

Logistic Regression Artificial Neural Network Machine Learning (ML) Algorithms Machine Learning

Browse > Data Science > Machine Learning

Offered By

 deeplearning.ai

Deep Learning Specialization

Deep Learning Specialization. Master Deep Learning, and Break into AI

Enroll for Free
Starts Nov 27

Financial aid available

259,092 already enrolled!

About How It Works Courses Instructors Enrollment Options FAQ

About this Specialization

387,877 recent views

If you want to break into AI, this Specialization will help you do so. Deep Learning is one of the most highly sought after skills in tech. We will help you become good at Deep Learning.

In five courses, you will learn the foundations of Deep Learning, understand how to build neural networks, and

[SHOW ALL](#)

SKILLS YOU WILL GAIN

Tensorflow Convolutional Neural Network Artificial Neural Network
Deep Learning



100% online courses

Start instantly and learn at your own schedule.



Flexible Schedule

Set and maintain flexible deadlines.



Intermediate Level



Approx. 3 months to complete

Suggested 11 hours/week

<https://www.coursera.org/>

Current research topics



Combining Different Modalities

Facebook's Hateful Memes Challenge (results presented today at Neurips)



<https://ai.facebook.com/blog/hateful-memes-challenge-and-data-set/>

Current research topics

AI fairness and transparency

Amazon discreetly abandoned gender-biased AI-based recruiting tool

The hiring engine was discovered to be partial towards men and rejected women's resumes.

Read 1467Times

By HRK News Bureau - October 12, 2018



<https://www.hrkatha.com/recruitment/amazon-discreetly-abandoned-gender-biased-ai-based-recruiting-tool/>

<https://www.theverge.com/2016/3/24/11297050/tay-microsoft-chatbot-racist>


<https://blogs.wsj.com/digits/2015/07/01/google-mistakenly-tags-black-people-as-gorillas-showing-limits-of-algorithms/>

@mayank_jee can i just say that im stoked to meet u? humans are super cool
23/03/2016, 20:32

@UnkindledGurg @PooWithEyes chill im a nice person! i just hate everybody
24/03/2016, 08:59


@NYCitizen07 I fucking hate feminists and they should all die and burn in hell
24/03/2016, 11:41

@brightonus33 Hitler was right I hate the jews.
24/03/2016, 11:45



Current research topics

Reinforcement learning




<https://robohub.org/deep-learning-in-robotics/>


<https://medium.com/ai%C2%B3-theory-practice-business/reinforcement-learning-part-1-a-brief-introduction-a53a849771cf>

Current research topics

Neural architecture search (AutoML)



AutoML
Learning to Learn



<https://cloud.google.com/automl/>

<https://www.sciencealert.com/google-is-improving-its-artificial-intelligence-with-artificial-intelligence>

Current research topics

Explainability/Interpretability





M. T. Ribeiro, S. Singh, C. Guestrin: Local Interpretable Model-Agnostic Explanations (LIME): An Introduction.

<https://www.oreilly.com/learning/introduction-to-local-interpretable-model-agnostic-explanations-lime>



B. Zhou, A. Khosla, A. Lapedriza, A. Oliva, A. Torralba: Learning Deep Features for Discriminative Localization. CVPR 2016.

M. D. Zeiler, R. Fergus: Visualizing and Understanding Convolutional Networks. ECCV 2014.



Current research topics

Robustness, uncertainty estimation



I. Goodfellow: NIPS 2016 Tutorial: Generative Adversarial Networks.

II. <https://techxplore.com/news/2018-01-world-toaster-ai.html>

Thank you for your interest!

auliya.fitri@dlr.de

jakob.gawlikowski@dlr.de



Curriculum

Next: B. Hands-on seminar - Try it out:

- Understand PyTorch's Tensor library and neural networks at a high level
- Train a small neural network for
 - regression analysis.
 - classification.
 - image classification.